

Exact Bayesian Inference for Spike-and-Slab Priors in Takagi–Sugeno–Kang Fuzzy Systems with Approximately Calibrated Model-Averaged Prediction Intervals

Zhiren Xiao^[ORCID]^b

^aGuangdong University of Finance, Guangzhou, China

Abstract

Calibrated prediction intervals are essential for decision-making under uncertainty, and Bayesian model selection is a principled route to them for rule-based regressions with linear consequents. However, replacing posterior sampling with an analytical approximation—BIC-based inclusion probabilities, hard thresholding, and a Laplace predictive covariance—can collapse accuracy. In the implementation we study, coverage collapses on ill-conditioned design matrices. We introduce exact inference (up to finite Monte Carlo error) for spike-and-slab priors over Takagi–Sugeno–Kang (TSK) consequents. A block-Gibbs sampler computes the per-rule inclusion full conditional in $O(n(d+1)^2)$ time via matrix-determinant and Woodbury identities; Bayesian model averaging (BMA) then propagates model uncertainty into the predictive variance. On two UCI benchmarks (three targets), this yields near-nominal 95% intervals (PICP = 0.93–0.95) where the analytical approximation we study undercovers (PICP = 0.00–0.18). Interval widths are comparable to the closed-form conjugate baseline’s on the Energy targets. A corrected dense baseline—a membership-function fix raising R^2 from 0.48 to 0.94 on the Energy-Cooling target—anchors the comparison. The approach applies to any rule-based regression with linear consequents, and the code and corrected baselines are released for reproducible, uncertainty-aware inference.

Keywords: uncertainty quantification, prediction intervals, Bayesian inference, spike-and-slab prior, Bayesian model averaging, Takagi–Sugeno–Kang fuzzy system

Exact Bayesian Inference for Spike-and-Slab Priors in Takagi–Sugeno–Kang Fuzzy Systems with Approximately Calibrated Model-Averaged Prediction Intervals

Zhiren Xiao^[ORCID]^b

^bGuangdong University of Finance, Guangzhou, China

1. Introduction

Predictive models in engineering increasingly need more than point estimates: designers require calibrated uncertainty to support decisions under limited data. Rule-based regressions with linear consequents—of which Takagi–Sugeno–Kang (TSK) fuzzy systems [1] are a canonical instance—are attractive here: they combine nonlinear approximation with a compact, interpretable architecture. Their linear consequents also make Bayesian treatment especially natural, yielding posterior distributions, prediction intervals, and principled model selection. Building-energy simulation exemplifies the data regime that motivates this work: each simulation run is computationally expensive, yielding training sets of only hundreds of samples [2, 3].

The Bayesian TSK literature has been hampered by two problems. The first is methodological. Spike-and-slab priors [4, 5] place positive probability on exact zero coefficients. A natural alternative to Markov chain Monte Carlo (MCMC) is analytical model selection—BIC-based inclusion probabilities followed by hard thresholding and a Laplace approximation to the predictive covariance. This approximation fails in two ways. Hard thresholding at posterior inclusion probability 0.5 discards rules that carry real (if partial) signal and collapses accuracy. The pipeline as commonly implemented—a ridge-regularized fit of the ill-conditioned design matrix combined with a Laplace predictive covariance on the pruned model—produces prediction intervals that cover only 0–18% of test points at a nominal 95% level. A faithful version of thresholding plus the Laplace covariance, fit by ordinary least squares, covers near-nominally but remains inaccurate and inefficient (Supplementary, Section 5). The approximation thus sacrifices accuracy, calibration, or both in exchange for a sparsity that, on the relevant benchmark, is not warranted.

The second problem is a *reproducibility* one: a subtle implementation error in the construction of fuzzy membership functions propagates through every downstream baseline. Specifically, an implementation pattern that computes the spreads of Gaussian membership functions from the training data during model fitting, but then *recomputes* them from query data (with incorrect cluster labels) at prediction time, silently depresses dense TSK baselines. The fuzzy partition at

prediction time no longer matches the one learned during training. We reproduce this pattern, quantify its effect, and show that correcting the bug raises the dense TSK least-squares baseline on the Energy Efficiency cooling target from $R^2 = 0.48$ to 0.94 (Section 5.1), providing a concrete corrected benchmark against which the methodological contribution below is validated.

We replace the analytical approximation with exact MCMC inference. We implement a block-Gibbs sampler for the spike-and-slab prior and predict by Bayesian model averaging (BMA), in which the predictive distribution is a mixture over posterior draws rather than a point estimate from a single selected model. Averaging over rule configurations retains partial information from uncertain rules and carries model uncertainty into the predictive variance, which avoids both the accuracy collapse and the interval collapse. We show that this recovers near-nominal prediction intervals, $\text{PICP} = 0.93\text{--}0.95$ at the nominal 95% level, where the analytical approximation it replaces undercovers severely as implemented ($\text{PICP} = 0.00\text{--}0.18$). Interval widths are comparable to the closed-form conjugate baseline’s on the Energy targets. To our knowledge, no prior work provides exact posterior sampling for spike-and-slab priors over rule-model consequents together with model-averaged prediction intervals.

Our contributions are threefold:

1. **Exact spike-and-slab inference for linear-consequent rule models.** A block-Gibbs sampler whose per-rule inclusion full conditional is computed in $O(n(d+1)^2)$ time, with Bayesian model averaging (BMA) propagating model uncertainty into the predictive variance.
2. **Calibrated model-averaged prediction intervals and a diagnostic of analytical approximations.** Gibbs sampling with BMA recovers near-nominal 95% intervals ($\text{PICP} = 0.93\text{--}0.95$), in contrast to our implementation of the BIC-plus-Laplace pipeline, which undercovers ($\text{PICP} = 0.00\text{--}0.18$); an ablation isolates the failure, showing that hard thresholding collapses accuracy while the as-implemented pipeline’s ridge-regularized fit of the ill-conditioned design matrix collapses coverage.
3. **When Sparse Selection Is Warranted.** We delineate when sparse selection is warranted: on low-dimensional data ($d = 8$) neither rule-level nor coefficient-level sparsity improves on the dense baseline. A real high-dimensional check ($d = 81$, Supplementary Table 7) reproduces the pattern rather than overturning it. The boundary is therefore not a low-dimensional artifact, and sparsity claims in this setting must be demonstrated rather than

86 assumed.

87 Section 2 reviews related work. Section 3 presents the TSK model, the membership bug, and
88 the Gibbs-based Bayesian estimators. Section 4 describes the experimental setup. Section 5 reports
89 the corrected baseline, the main comparison, calibration, and the sparsity boundary. Section 6
90 discusses implications and limitations, and Section 7 concludes.

91 2. Related Work

92 2.1. Bayesian TSK Fuzzy Systems

93 Gu and colleagues introduced Bayesian TSK classifiers with conjugate Gaussian–Wishart
94 priors [6], later extended to joint structure identification [7], imbalanced classification [8], and
95 probabilistic Gaussian centralization [9]. A zero-order Bayesian TSK variant is studied in
96 [10], prediction intervals for adaptive neuro-fuzzy inference systems have been constructed
97 through interval-network training in [11], evolving fuzzy systems maintain interval predictions
98 in nonstationary environments [12], and interval-construction methods for neural predictors are
99 surveyed in [13]. None of these targets sparse high-order consequents through posterior model
100 averaging. These frameworks provide full posterior uncertainty estimates (posterior variances
101 and prediction intervals) but employ non-sparsity-inducing priors, so every rule remains active
102 regardless of relevance [14]. Our work builds directly on this line: the dense conjugate posterior
103 (Section 3.3) is the natural non-sparsity-inducing reference against which sparse Bayesian variants
104 should be judged, and we use it to validate our MCMC implementation.

105 2.2. Sparse TSK and Bayesian Variable Selection

106 On the frequentist side, Gong et al. [15] combine LASSO with TSK consequents for embedded
107 feature selection, and Zhang et al. [16] use group-LASSO regularization for nonstationary TSK
108 networks; convergence guarantees are studied in [17]. These methods yield point estimates without
109 uncertainty quantification. Exact posterior computation for Bayesian model selection is generally
110 intractable, and analytical surrogates built on information criteria or Laplace approximations have
111 well-documented limitations [18, 19]. Recent work makes the error of Laplace approximations
112 quantitative and finitely computable [20]. On the Bayesian side, spike-and-slab priors [4, 5, 21]
113 have been applied to structurally sparse Bayesian neural networks via variational inference [22]
114 and to high-dimensional sensor selection in engineering prognostics [23]. Denœux’s ENNreg [24]
115 provides a belief-function alternative to probabilistic fuzzy regression with prediction intervals

but without rule selection. High-dimensional TSK systems target regimes with far more features: Bian et al. [25] introduce a multi-head TSK system for high-dimensional data, and Xue et al. [26] combine feature selection with rule extraction. Our low-dimensional benchmarks ($d = 8$) do not exercise this regime; we return to this contrast in the sparsity-boundary discussion (Section 5.4). Distribution-free alternatives such as *split* conformal prediction [27, 28] provide finite-sample coverage without posterior sampling. It requires a separate calibration set and does not decompose predictive variance into noise and model-selection components, which our BMA construction propagates explicitly. Our contribution is to implement exact (MCMC) spike-and-slab inference for TSK consequents with BMA-calibrated prediction intervals, which to our knowledge has not been developed for linear-consequent rule models, and to evaluate it against the dense baseline.

2.3. Reproducibility in Soft Computing

Reproducibility failures arising from silent implementation errors are a recognized concern in applied machine learning [29]. Within fuzzy systems, the practical consequences of an inconsistent fuzzy partition between training and inference have not, to our knowledge, been quantified. We quantify them on a reproduction of this pattern and release the corrected code and baselines publicly.

3. Methodology

3.1. TSK Fuzzy System

For input $\mathbf{x} \in \mathbb{R}^d$ and scalar target y , a TSK system with R rules is

$$y = \sum_{j=1}^R w_j(\mathbf{x}) (\beta_{j0} + \beta_{j1}x_1 + \cdots + \beta_{jd}x_d) + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2), \quad (1)$$

where $w_j(\mathbf{x}) = \prod_{i=1}^d \mu_{A_{ji}}(x_i) / \sum_{k=1}^R \prod_{i=1}^d \mu_{A_{ki}}(x_i)$ is the normalized firing strength of rule j , and A_{ji} are Gaussian antecedent fuzzy sets.

$$\mu_{A_{ji}}(x_i) = \exp\left(-\frac{(x_i - c_{ji})^2}{2s_{ji}^2}\right), \quad (2)$$

Centers c_{ji} are obtained from fuzzy c -means clustering [30], and spreads s_{ji} are set to the standard deviation of the training points assigned to each cluster, scaled by a factor s . Writing the weighted design matrix $\Phi \in \mathbb{R}^{n \times R(d+1)}$ and stacking consequents into $\beta \in \mathbb{R}^{R(d+1)}$, Eq. (1) becomes the linear model $\mathbf{y} = \Phi\beta + \varepsilon$.

141 3.2. The Membership-Spread Bug

142 The spreads s_{ji} in Eq. (2) define the fuzzy partition and must be identical at training and
 143 inference time. A defect we reproduce recomputes them at prediction: it calls the membership
 144 routine on query inputs while passing a placeholder label vector (e.g., all zeros), so the routine
 145 re-estimates spreads from the query data with degenerate cluster assignments. The resulting
 146 partition no longer matches the one used to fit β . In our reproduction, this single defect depresses
 147 the dense TSK baseline from $R^2 = 0.94$ to 0.48 on the Energy-Cooling target (Section 5.1).
 148 The correct procedure stores the training-time centers and spreads and reuses them verbatim at
 149 inference, as in Algorithm 1.

Algorithm 1 Corrected TSK training and inference

Require: Training data $\{\mathbf{x}_i, y_i\}_{i=1}^n$, rule count R

- 1: $(\mathbf{c}, \mathbf{u}) \leftarrow \text{FCM}(\mathbf{X}, R)$ ▷ fuzzy c -means
 - 2: $\mathbf{s} \leftarrow \{\text{per-cluster standard deviations from training points, scaled by } s\}$
 - 3: Store $(\mathbf{c}, \mathbf{s})$ ▷ fix the partition at training time
 - 4: Build Φ via Eqs. (1)–(2) using $(\mathbf{c}, \mathbf{s})$
 - 5: Fit β (least squares, conjugate Bayes, or Gibbs; Sections 3.3–3.4)
 - 6: **return** model with frozen $(\mathbf{c}, \mathbf{s}, \beta)$
-

150 3.3. Conjugate Bayesian Baseline

151 The non-sparse Bayesian TSK places a Gaussian prior $\beta \sim \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I})$ and an inverse-gamma
 152 prior $\sigma^2 \sim \text{InvGamma}(0.01, 0.01)$ on the noise variance. The posterior is available in closed form:
 153 with $\mathbf{P} = \Phi^\top \Phi + \tau^{-2} \mathbf{I}$,

$$\beta \mid \mathcal{D}, \sigma^2 \sim \mathcal{N}(\mathbf{P}^{-1} \Phi^\top \mathbf{y}, \sigma^2 \mathbf{P}^{-1}), \quad (3)$$

154 and the predictive variance at a new design row ϕ^* is $\sigma^2 + \phi^* (\sigma^2 \mathbf{P}^{-1}) \phi^{*\top}$. This closed-form model
 155 serves both as a calibrated reference and as a correctness check for our sampler (a Gibbs sampler
 156 with all rules forced active must recover it). When the block-Gibbs sampler of Section 3.4 is
 157 run with all inclusion indicators fixed to $\gamma_j = 1$, the sampled predictive distribution recovers
 158 the conjugate closed form on held-out data. The predictive mean matches in R^2 to within 10^{-3}
 159 (measured 0.0010 at the default $R = 5$), and the predictive variance agrees to within $\sim 6\%$ (median;
 160 measured 5.6%). These results confirm the correctness of the β and σ^2 updates against the closed
 161 form (reproduction log in `results/raw/gibbs_verify_output.txt`). The inclusion-indicator
 162 full conditional of Eq. (5) is not exercised by this fixed-activation check. We validate it separately
 163 on a synthetic problem with known sparse ground truth, where the sampler’s inclusion probabilities

164 agree with exhaustive model averaging to within 10^{-3} (Supplementary, Section 9). We compare at
 165 the predictive level: the normal-equations matrix $\Phi^\top \Phi$ is numerically ill-conditioned (condition
 166 number $\sim 10^{18}$), which makes the raw posterior covariance numerically unstable for both the
 167 closed form and any MCMC estimate. The regularized estimation problem is nevertheless well
 168 posed. The dense fits use ridge-regularized normal equations (penalty 10^{-4} on $\Phi^\top \Phi$), which
 169 stabilizes them.

170 3.4. Spike-and-Slab Prior and Gibbs Sampling

171 We decompose β into R blocks $\beta_j \in \mathbb{R}^{d+1}$ and place a spike-and-slab prior on each block:

$$\beta_j \mid \gamma_j, \tau^2 \sim \begin{cases} \mathcal{N}(\mathbf{0}, \tau^2 \mathbf{I}_{d+1}), & \gamma_j = 1 \text{ (slab)}, \\ \delta_{\mathbf{0}}, & \gamma_j = 0 \text{ (spike)}, \end{cases} \quad (4)$$

172 with $\gamma_j \sim \text{Bernoulli}(\pi)$. Rather than approximate the posterior over γ with a BIC criterion, we
 173 sample it exactly with a block-Gibbs sampler. Here “exact” means that we target the posterior by
 174 sampling rather than replacing it with a BIC-plus-Laplace analytical surrogate; as with any MCMC
 175 procedure the estimates carry finite Monte Carlo error, which we quantify with the convergence
 176 diagnostics of Section 4. For block j , let $\mathbf{r}_{-j} = \mathbf{y} - \sum_{k \neq j} \Phi_k \beta_k$ be the residual excluding rule j ,
 177 and let $\mathbf{A}_j = \Phi_j^\top \Phi_j$. The marginal likelihood of $\mathbf{r}_{-j}$ under $\gamma_j = 1$ (integrating out β_j) is Gaussian
 178 with covariance $\sigma^2 \mathbf{I} + \tau^2 \Phi_j \Phi_j^\top$. Its determinant and quadratic form are computed in $O(n(d+1)^2)$
 179 via the matrix-determinant and Woodbury identities. The full conditional for γ_j is

$$P(\gamma_j = 1 \mid \cdot) = \sigma\left(\log \frac{\pi}{1-\pi} - \frac{1}{2} \log \det(\mathbf{I} + \frac{\tau^2}{\sigma^2} \mathbf{A}_j) + \frac{\tau^2}{2\sigma^4} \mathbf{z}_j^\top (\mathbf{I} + \frac{\tau^2}{\sigma^2} \mathbf{A}_j)^{-1} \mathbf{z}_j\right), \quad (5)$$

180 where $\mathbf{z}_j = \Phi_j^\top \mathbf{r}_{-j}$ and $\sigma(\cdot)$ is the logistic function. Conditional on $\gamma_j = 1$, we draw $\beta_j \sim$
 181 $\mathcal{N}((\mathbf{A}_j/\sigma^2 + \tau^{-2} \mathbf{I})^{-1} \mathbf{z}_j/\sigma^2, (\mathbf{A}_j/\sigma^2 + \tau^{-2} \mathbf{I})^{-1})$; otherwise $\beta_j = \mathbf{0}$. Finally σ^2 is drawn from its
 182 inverse-gamma full conditional. A coefficient-level variant applies the same prior independently
 183 to each scalar coefficient β_i (stochastic search variable selection, SSVS [5]), which is appropriate
 184 when feature-level rather than rule-level sparsity is of interest. For numerical stability on the
 185 ill-conditioned design matrix, the implementation adds a small diagonal jitter (10^{-10}) to the
 186 sampled precision matrices, initializes the coefficients from a ridge-regularized least-squares fit,
 187 and floors the noise variance at 10^{-4} .

188 3.5. Bayesian Model Averaging for Prediction

189 We predict by BMA rather than by conditioning on a single selected model [31]; recent
 190 reviews of BMA practice and its pitfalls are given in [32]. Given T posterior draws $\{\boldsymbol{\beta}^{(t)}, \sigma^{2(t)}\}_{t=1}^T$,
 191 the predictive mean is the mixture mean $\hat{y}^* = \frac{1}{T} \sum_t \boldsymbol{\phi}^* \boldsymbol{\beta}^{(t)}$, and the predictive variance follows
 192 from the law of total variance:

$$\text{Var}(y^* | \mathbf{x}^*, \mathcal{D}) = \underbrace{\frac{1}{T} \sum_t \sigma^{2(t)}}_{\text{aleatoric}} + \underbrace{\text{Var}_t(\boldsymbol{\phi}^* \boldsymbol{\beta}^{(t)})}_{\text{model uncertainty}}, \quad (6)$$

193 yielding the 95% interval $\hat{y}^* \pm 1.96 \sqrt{\text{Var}(y^* | \mathbf{x}^*, \mathcal{D})}$. Over posterior draws the second term spans
 194 both between-model (rule-configuration) variation and within-model parameter variation. It is
 195 precisely the model-selection uncertainty that a hard-thresholded approximation discards, and
 196 retaining it is what restores calibration.

197 4. Experimental Setup

198 We evaluate on two UCI regression benchmarks (three targets): the Energy Efficiency dataset
 199 [2] (768 simulated building configurations, 8 geometric features, with heating and cooling loads
 200 as two targets) and the Concrete Compressive Strength dataset [33] (1030 laboratory samples, 8
 201 mix-design features). Features are standardized to zero mean and unit variance; all TSK methods
 202 use $R = 5$ rules from fuzzy c -means ($m = 2$) with Gaussian membership spreads equal to the
 203 per-cluster standard deviation of training points, scaled by $s = 1.5$. A rule-count sweep ($R = 3$ –10,
 204 Supplementary) shows dense accuracy is nearly unchanged from $R = 4$ upward on the Energy
 205 targets and improves only marginally beyond $R \approx 7$ on Concrete, supporting the default $R = 5$.

206 Four TSK methods are compared in the main experiments: **TSK-LS** (dense least squares),
 207 **Bayesian-TSK** (dense conjugate posterior, Section 3.3), **TSK-SpikeSlab-BIC** (the BIC-plus-
 208 Laplace analytical approximation with hard thresholding; labeled SpikeSlab-Fast in the re-
 209 leased code), and **TSK-SpikeSlab-Gibbs** (rule-level Gibbs plus BMA, Sections 3.4–3.5; labeled
 210 SpikeSlab-Gibbs). A coefficient-level variant, **TSK-SSVS** (labeled SSVS-Gibbs), is used in
 211 the sparsity-boundary analysis (Section 5.4). Random forest, support vector regression, and
 212 Gaussian-process regression serve as non-interpretable machine-learning references; the Gaussian
 213 process provides closed-form predictive intervals and is therefore the key probabilistic baseline for
 214 the calibration comparison. Reference implementations use the scikit-learn configurations shipped

215 in the released code. A random forest with 300 trees and maximum depth 10, support vector
 216 regression with an RBF kernel ($C = 1$, $\gamma = \text{scale}$), and a Gaussian process with an RBF-plus-
 217 white-noise kernel, hyperparameters optimized by marginal likelihood (`normalize_y=True`).
 218 Fuzzy-specific interval-construction methods such as ENNreg [24] and interval-ANFIS [11]
 219 build prediction intervals through evidential combination or interval-network training rather than
 220 posterior model selection. They are therefore not direct baselines for the model-selection question
 221 addressed here, and the Gaussian process—whose closed-form intervals come from a standard
 222 library implementation—serves as the probabilistic reference. The rule-level Gibbs sampler uses
 223 1,000 burn-in and 2,000 retained draws; the coefficient-level SSVS uses 800 and 800. For the
 224 rule-level sampler, convergence is verified by running three chains with independent seeds on
 225 representative splits. The Gelman–Rubin statistic satisfies $\hat{R} \leq 1.001$ for the consequents and the
 226 noise variance, and the effective sample size exceeds 1,700 for the noise variance (Supplementary
 227 Table 6). The inclusion-indicator chains saturate near full activation on these benchmarks (Sec-
 228 tion 5.4), so their effective sample size equals the chain length by construction rather than from
 229 active exploration. Formal convergence verification is claimed only for the rule-level sampler,
 230 not for the coefficient-level SSVS runs. All results are averaged over 30 independent random
 231 80/20 train–test splits with a fixed seed (SEED=42); reported intervals are 95% Gaussian intervals.
 232 All spike-and-slab priors use inclusion probability $\pi = 0.5$; the rule-level Gibbs and conjugate
 233 baselines use slab variance $\tau^2 = 10^3$, and the coefficient-level SSVS uses $\tau^2 = 1.0$. The $\pi = 0.5$
 234 prior is dominated by the data on these benchmarks (posterior inclusion probabilities saturate near
 235 one), and the rule-level slab $\tau^2 = 10^3$ is deliberately diffuse relative to the standardized-feature
 236 scale. The sensitivity of the coefficient-level variant to τ^2 is characterized in Fig. 6a. Metrics are
 237 root-mean-square error (RMSE), R^2 , prediction-interval coverage probability (PICP), and mean
 238 prediction-interval width (MPIW).

239 5. Results

240 5.1. A Corrected Dense Baseline

241 Fig. 1 contrasts the R^2 values reproduced under the buggy membership implementation with
 242 those obtained after correcting it. The improvement is consistent: TSK-LS rises from 0.48 to 0.94
 243 (Energy-Cooling), 0.78 to 0.97 (Energy-Heating), and -2.86 to 0.74 (Concrete). The conjugate
 244 Bayesian-TSK rises from 0.78 to 0.94, 0.84 to 0.97, and -0.56 to 0.74. This confirms that the

bug, not the model class, was responsible for the apparently mediocre baselines, and it establishes
the corrected numbers as the reference for the comparisons below.




Figure 1: Corrected dense baselines. Reproduced R^2 under the buggy membership implementation (orange) versus corrected R^2 (blue, training-time membership spreads reused at inference) for the dense TSK-LS and conjugate Bayesian-TSK baselines. Values are means over 30 splits on the Energy (Heating and Cooling) and Concrete targets; correcting the membership-spread bug raises the dense TSK-LS R^2 from 0.48 to 0.94 on Energy-Cooling.

5.2. Main Comparison

Table 1 and Fig. 2 report the primary results (full mean-absolute-error results in Supplementary Table 1). After correcting the bug, the dense baselines are strong: TSK-LS achieves $R^2 = 0.94$ – 0.97 on the Energy targets and 0.74 on Concrete, and the conjugate Bayesian-TSK differs from it by at most 0.01 in R^2 . The proposed Gibbs sampler recovers this accuracy: TSK-SpikeSlab-Gibbs attains $R^2 = 0.97/0.94/0.74$, matching the dense baseline to within 0.01 in R^2 on every target. This is the expected behavior of a correct sampler and a direct validation of our MCMC implementation against the closed form. A paired Wilcoxon signed-rank test over the 30 splits [34] finds no significant difference between TSK-SpikeSlab-Gibbs and TSK-LS on the Energy targets (Heating $p = 0.11$, Cooling $p = 0.07$). It also finds no significant difference between Bayesian-TSK and TSK-LS (Heating $p = 0.32$, Cooling $p = 0.87$). The test does not establish equivalence, and the split-level results are not independent because the splits reuse the same underlying rows. On Concrete the difference is nominally significant ($p < 0.01$) but tiny in magnitude, with the samplers slightly above the dense baseline (mean $\Delta R^2 = 0.006$ for Gibbs, 0.004 for Bayesian-TSK).

By contrast, the BIC-based analytical approximation TSK-SpikeSlab-BIC collapses to negative R^2 (-1.8 to -6.0), significantly worse than every other method on all three targets (all $p < 0.001$). It predicts worse than a constant mean because hard thresholding discards rules that the correctly specified model needs.

The Gaussian-process reference [35] attains the accuracy of the non-interpretable Random Forest (GP $R^2 = 0.998/0.98/0.89$). Its closed-form 95% intervals slightly under-cover (PICP = 0.93 – 0.94), whereas the exact Bayesian TSK methods reach near-nominal coverage (PICP = 0.93 – 0.95) at the cost of wider intervals. On the Heating and Concrete targets the TSK intervals sit closer to the nominal level than the GP’s, while on Cooling the two are nearly tied (TSK 0.93 vs GP 0.94).

Table 1: Main comparison (mean over 30 splits; $\pm$ std shown for RMSE; R^2 and PICP are means). PICP is the 95% interval coverage. Differences among TSK-LS, Bayesian-TSK, and TSK-SpikeSlab-Gibbs are at most 0.01 in R^2 ; paired Wilcoxon tests over the 30 splits are reported in the text.

Method	Energy-Heating			Energy-Cooling			Concrete		
	RMSE	R^2	PICP	RMSE	R^2	PICP	RMSE	R^2	PICP
TSK-LS	1.69 $\pm$ 0.19	0.97	–	2.36 $\pm$ 0.17	0.94	–	8.52 $\pm$ 0.55	0.74	–
Bayesian-TSK	1.69 $\pm$ 0.19	0.97	0.95	2.36 $\pm$ 0.17	0.94	0.93	8.46 $\pm$ 0.43	0.74	0.95
TSK-SpikeSlab-BIC	16.4 $\pm$ 3.4	–1.8	0.00	18.8 $\pm$ 3.0	–3.1	0.00	44.0 $\pm$ 4.0	–6.0	0.18
TSK-SpikeSlab-Gibbs	1.70 $\pm$ 0.19	0.97	0.95	2.37 $\pm$ 0.17	0.94	0.93	8.43 $\pm$ 0.43	0.74	0.95
Random Forest	0.49 $\pm$ 0.06	0.998	–	1.67 $\pm$ 0.13	0.97	–	5.17 $\pm$ 0.40	0.90	–
SVR	2.73 $\pm$ 0.21	0.93	–	3.09 $\pm$ 0.22	0.89	–	9.97 $\pm$ 0.43	0.64	–
Gaussian Process	0.47 $\pm$ 0.05	0.998	0.93	1.35 $\pm$ 0.11	0.98	0.94	5.53 $\pm$ 0.46	0.89	0.94




Figure 2: Predictive accuracy (R^2) across datasets. The Gibbs/BMA sampler (TSK-SpikeSlab-Gibbs) recovers the dense-baseline accuracy on all three targets (Energy-Heating, Energy-Cooling, Concrete; $R^2 = 0.97/0.94/0.74$), while the BIC-based analytical approximation (TSK-SpikeSlab-BIC) collapses to negative R^2 .

272 5.3. Calibration

273 Fig. 3 reports interval calibration, assessed with the coverage and width diagnostics standard
 274 in the probabilistic-forecast literature [36]. Pooled over the 30 splits, both the conjugate Bayesian-
 275 TSK and the Gibbs sampler attain mean PICP of 0.95 on the Energy-Heating and Concrete targets
 276 and 0.935 on Energy-Cooling (displayed as 0.93 in Table 1). The 95% binomial (Wilson) intervals
 277 are [0.945, 0.958], [0.944, 0.955], and [0.927, 0.941], respectively. The Gibbs values agree with
 278 the conjugate values to within 0.004. Coverage is therefore within one binomial standard error of
 279 the nominal 95% on Heating and Concrete, and about 1.5 percentage points below it on Cooling.
 280 The deviation lies within one per-split standard deviation (~ 0.027 on Cooling) yet is statistically
 281 detectable once the splits are pooled. We therefore describe the intervals as near-nominal rather
 282 than perfectly calibrated. Widths are compact ($\text{MPIW} \approx 6.7\text{--}32.9$; Supplementary Table 8). The
 283 analytical TSK-SpikeSlab-BIC baseline undercovers severely as implemented: $\text{PICP} = 0.00$
 284 on both Energy targets and 0.18 on Concrete. All but a small fraction of test points on the
 285 Energy targets, and the large majority on Concrete, fall outside the nominal 95% intervals. The
 286 ablation (Supplementary, Section 5) isolates the mechanism: hard thresholding alone collapses
 287 accuracy ($R^2 = 0.60$) while retaining near-nominal coverage with inflated width ($\text{MPIW} 26.1$).
 288 The coverage collapse thus arises from the as-implemented pipeline’s ridge-regularized fit of the
 289 ill-conditioned design matrix rather than from the Laplace covariance per se. The Gibbs/BMA
 290 prediction of Eq. (6) restores accuracy and efficiency, propagating model uncertainty through
 291 the mixture variance. The value of averaging over configurations is most visible where the
 292 closed-form conjugate baseline degrades. On the high-dimensional check (Superconductivity,
 293 $d = 81$, Section 5.4, Supplementary Table 7), the Gibbs/BMA estimator maintains near-nominal
 294 calibration ($\text{PICP} 0.940$) where the dense conjugate baseline slightly undercovers ($\text{PICP} 0.921$).
 295 Fig. 4 confirms this calibration across a range of nominal levels, and Fig. 5 shows the 95%
 296 intervals of the two Bayesian methods against the held-out truth on a representative split.




Figure 3: Prediction-interval coverage (PICP) with the nominal 95% reference (dashed). Gibbs/BMA and the conjugate Bayesian model attain near-nominal coverage; the BIC-plus-Laplace approximation undercovers as implemented.




Figure 4: Empirical coverage versus nominal coverage for the conjugate Bayesian-TSK, the Gibbs/BMA sampler, and the Gaussian-process reference on Energy-Cooling (30 splits). Points on the diagonal would indicate perfectly calibrated intervals. The Bayesian methods over-cover at the lowest nominal levels (empirical coverage 0.57 at nominal 0.5). They under-cover slightly in the range 0.8–0.95 (empirical 0.94 at nominal 0.95) and return to the diagonal at nominal 0.99. The Gaussian process falls below the diagonal at nominal 0.99.




Figure 5: 95% prediction intervals (shaded) and mean prediction (line) versus held-out truth (points) on one representative Energy-Cooling split, for the conjugate Bayesian-TSK (left) and Gibbs+BMA (right). The intervals track the sorted truth closely and cover the large majority of test points at compact widths.

297 5.4. The Sparsity Boundary

298 Having established that Gibbs/BMA is correct and calibrated, we characterize when the
 299 *sparsity* component of the spike-and-slab prior helps. Two results stand out (Fig. 6).

300 First, rule-level sparsity rarely activates: on all three targets the posterior inclusion probabilities
 301 of the $R = 5$ rules are close to 1.0 (with a single rule occasionally pruned on the Energy targets),
 302 because with a correct partition nearly every rule contributes. This is expected in the low-
 303 dimensional regime ($d = 8$), where the dense estimation problem is well posed. The rule-count
 304 sweep (Supplementary) confirms that even at $R = 10$ the inclusion probabilities remain near 1.0,
 305 so low dimensionality, not the rule count, keeps rule-level sparsity inactive.

306 Second, coefficient-level sparsity (SSVS) is a trade-off, not an accuracy gain. Varying the slab
 307 variance τ^2 over three orders of magnitude (Fig. 6a) shows that a tight slab ($\tau^2 = 0.1$) over-shrinks
 308 and collapses accuracy to $R^2 = -5.4$. A diffuse slab ($\tau^2 = 100$), in contrast, merely recovers the
 309 dense baseline ($R^2 = 0.92$ versus 0.94) while pruning nothing. No τ^2 setting improves on the
 310 dense baseline. In the irrelevant-feature regime (Fig. 6b), appending up to 30 standard-normal
 311 noise features leaves the dense baseline near its original level (ending at $R^2 = 0.88$, with a small
 312 non-monotone rise at four features that is within split-level noise). The sparse variants do not
 313 improve upon it. Correct fuzzy c -means already provides robustness to irrelevant features, so
 314 sparsity does not improve on the dense baseline in this low-dimensional regime.

Sparsity boundary under correct fuzzy c-means




Figure 6: The sparsity boundary on Energy-Cooling ($R = 5$). (a) Coefficient-level SSVS R^2 versus slab variance τ^2 (10 splits): a tight slab over-shrinks and a diffuse slab only recovers the dense baseline (dashed; 30-split TSK-LS reference, $R^2 = 0.94$), never improving on it. (b) Irrelevant-feature regime (10 splits; dense baseline $R^2 = 0.93$ at zero noise features): the dense TSK-LS declines only modestly, from $R^2 = 0.93$ to 0.88, as noise features are appended. A small non-monotone rise at four features is within split-level noise. Neither rule-level nor coefficient-level sparsity improves on it (details in Supplementary Tables 2–3).

6. Discussion

6.1. Why the Analytical Approximation Fails

The failure of the BIC-plus-Laplace approximation has a specific origin. BIC approximates the Bayes factor [37] by replacing the integral over the slab with a point mass at the maximum-likelihood estimate. This is only reliable when the posterior over γ_j is concentrated at 0 or 1. On the Energy benchmark the inclusion probabilities are intermediate, so the hard threshold at 0.5 discards rules that carry real (if partial) signal. The Laplace predictive covariance then conditions on a single pruned model. With an ordinary-least-squares fit it alone yields wide intervals with near-nominal coverage (MPIW 26.1, PICP 0.93). The coverage collapse instead arises in the as-implemented pipeline, which fits the selected model with ridge regularization on the ill-conditioned design matrix. Gibbs sampling with BMA avoids both steps: it averages over all rule configurations rather than selecting one. An ablation isolates which step is decisive (Supplementary, Section 5): replacing hard thresholding with BMA over all 2^R rule subsets, using the same BIC-based model weights, restores dense-level accuracy and near-nominal calibration ($R^2 = 0.95$, PICP 0.92 on Energy-Cooling). Averaging rather than the BIC approximation is the decisive change for accuracy. Feeding the Gibbs sampler’s inclusion probabilities through

331 hard thresholding and the Laplace covariance performs equally well ($R^2 = 0.94$, PICP 0.94),
 332 confirming that the Gibbs inclusion probabilities are reliable. The as-implemented BIC baseline
 333 collapses ($R^2 = -2.5$, PICP 0.00; 10-split ablation, Supplementary Section 5, cf. -3.1 at 30 splits
 334 in Table 1). The coverage collapse is traceable to its ridge-regularized fit of the ill-conditioned
 335 design matrix.

336 6.2. Implications for Sparse TSK Research

337 Our boundary characterization carries a constructive message. Sparsity-inducing priors are
 338 worth their cost only when the model is over-parameterized (high-dimensional inputs with truly
 339 irrelevant features) or when the data are scarce relative to the number of effective parameters. Our
 340 real high-dimensional check ($d = 81$, Supplementary Table 7) does not realize this regime, since
 341 its features are largely relevant and $n = 3,000$ is not scarce relative to the effective parameter
 342 count. In the low-dimensional building-energy regime, the dense TSK is already well posed and
 343 no sparsity is warranted. The rule-count sweep (Supplementary) sharpens this: even at $R = 10$,
 344 where the model has $R(d+1) = 90$ consequent coefficients, rule-level sparsity does not activate.
 345 Rule count alone, without high dimensionality, does not make sparse selection advantageous on
 346 these benchmarks. This is not a failure of the spike-and-slab framework; it is a statement about the
 347 data, and it pinpoints where the framework’s cost is justified. Alternative shrinkage priors, such
 348 as the continuous spike-and-slab LASSO [38] and the horseshoe [39], offer smooth regularizers
 349 that may behave differently under the collinearity of the TSK weighted design matrix. Extending
 350 the exact inference developed here to these priors is natural future work. We therefore present
 351 the Gibbs/BMA estimator as a *correct and calibrated* Bayesian TSK, and report sparsity as a
 352 boundary condition rather than a primary contribution.

353 6.3. Limitations

354 Three limitations scope this work. First, the benchmarks are low-dimensional ($d = 8$),
 355 although an evaluation on a real high-dimensional benchmark (Superconductivity, $d = 81$,
 356 subsample of $n = 3,000$ rows, Supplementary Table 7) reproduces the same pattern. Rule-level
 357 sparsity matches but does not improve on the dense baseline, and coefficient-level sparsity over-
 358 shrinks. The boundary is therefore not a low-dimensional artifact; whether sparsity is beneficial
 359 under different antecedent structures or much larger samples remains open. Second, we fix the
 360 antecedent structure and apply sparsity only to consequents; joint antecedent–consequent selection

[7] remains future work. Third, our evaluation uses simulated (Ecotect) building-energy data; validation against measured, submetered energy data would strengthen the application case.

7. Conclusion

We introduced exact Bayesian inference for spike-and-slab priors over TSK consequents. A block-Gibbs sampler combined with Bayesian model averaging recovers near-nominal 95% prediction intervals (PICP = 0.93–0.95) where our implementation of the BIC-plus-Laplace analytical pipeline undercovers (PICP = 0.00–0.18). In practical terms, a modeler who needs a 95% interval from a rule-based model should average over rule configurations rather than commit to a single selected model. The averaged intervals track the nominal level closely, whereas the thresholded BIC baseline undercovers. We validated the method against a corrected dense baseline, obtained by fixing a membership-function bug that had depressed TSK baselines ($R^2 = 0.48 \rightarrow 0.94$ on the Energy-Cooling target). We also characterized when sparsity helps, showing that on low-dimensional data neither rule-level nor coefficient-level sparsity improves on the dense baseline, delineating the regime in which sparse selection is warranted. The approach applies to any rule-based regression with linear consequents, and the corrected code and baselines are released to support reproducible, uncertainty-aware inference.

CRedit authorship contribution statement

Zhiren Xiao: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review & editing, Visualization.

Acknowledgments

The author declares no specific funding for this research.

Data Availability

Datasets are publicly available from the UCI Machine Learning Repository: Energy Efficiency (dataset 242) [2] and Concrete Compressive Strength (dataset 165) [33]. Code is available at <https://github.com/Jackxiaozhiren/tsk-spikeslab> and archived with the persistent DOI [10.5281/zenodo.21929319](https://doi.org/10.5281/zenodo.21929319), including the corrected TSK core, experiment drivers, a pinned dependency manifest, and the fixed seed stated in Section 4.

388 **Declaration of Competing Interests**

389 The author declares no competing interests.

390 **Declaration of Generative AI and AI-assisted technologies in the writing process**

391 During the preparation of this work the author used Claude (Anthropic) for code development
392 assistance and language editing. All scientific content, methodology design, experimental results,
393 and interpretation were performed and verified by the author. After using this tool, the author
394 reviewed and edited the content as needed and takes full responsibility for the content of the
395 published article.

396 **References**

- 397 [1] T. Takagi, M. Sugeno, Fuzzy identification of systems and its applications to modeling and
398 control, IEEE Transactions on Systems, Man, and Cybernetics SMC-15 (1) (1985) 116–132.
399 [doi:10.1109/TSMC.1985.6313399](https://doi.org/10.1109/TSMC.1985.6313399).
- 400 [2] A. Tsanas, A. Xifara, Accurate quantitative estimation of energy performance of residential
401 buildings using statistical machine learning tools, Energy and Buildings 49 (2012) 560–567.
402 [doi:10.1016/j.enbuild.2012.03.003](https://doi.org/10.1016/j.enbuild.2012.03.003).
- 403 [3] G. Guo, P. Liu, Y. Zheng, Early energy performance analysis of smart buildings by
404 consolidated artificial neural network paradigms, Heliyon 10 (4) (2024) e25848. [doi:
405 10.1016/j.heliyon.2024.e25848](https://doi.org/10.1016/j.heliyon.2024.e25848).
- 406 [4] T. J. Mitchell, J. J. Beauchamp, Bayesian variable selection in linear regression, Journal of
407 the American Statistical Association 83 (404) (1988) 1023–1032. [doi:10.2307/2290129](https://doi.org/10.2307/2290129).
- 408 [5] E. I. George, R. E. McCulloch, Variable selection via Gibbs sampling, Journal of the
409 American Statistical Association 88 (423) (1993) 881–889. [doi:10.2307/2290777](https://doi.org/10.2307/2290777).
- 410 [6] X. Gu, F.-L. Chung, S. Wang, Bayesian Takagi–Sugeno–Kang fuzzy classifier, IEEE Transac-
411 tions on Fuzzy Systems 25 (6) (2017) 1655–1671. [doi:10.1109/TFUZZ.2016.2617377](https://doi.org/10.1109/TFUZZ.2016.2617377).
- 412 [7] X. Gu, S. Wang, Bayesian Takagi–Sugeno–Kang fuzzy model and its joint learning of struc-
413 ture identification and parameter estimation, IEEE Transactions on Industrial Informatics
414 14 (12) (2018) 5327–5337. [doi:10.1109/TII.2018.2813977](https://doi.org/10.1109/TII.2018.2813977).

- [8] X. Gu, F.-L. Chung, H. Ishibuchi, S. Wang, Imbalanced TSK fuzzy classifier by cross-class Bayesian fuzzy clustering and imbalance learning, *IEEE Transactions on Systems, Man, and Cybernetics: Systems* 47 (8) (2017) 2005–2020. doi:10.1109/tsmc.2016.2598270.
- [9] Y. Li, F.-L. Chung, Y. Nojima, S. Wang, Fully interpretable Gaussian centralized TSK fuzzy classifier from probabilistic perspective: Concepts, output-stability-based learning, and ensemble, *IEEE Transactions on Fuzzy Systems* (2025). doi:10.1109/tfuzz.2025.3600911.
- [10] J. Liu, F.-L. Chung, S. Wang, Bayesian zero-order Takagi–Sugeno fuzzy system modeling, *Applied Soft Computing* 55 (2017) 140–151. doi:10.1016/j.asoc.2017.01.040.
- [11] B. Miskony, D. Wang, Construction of prediction intervals using adaptive neurofuzzy inference systems, *Applied Soft Computing* 72 (2018) 579–586. doi:10.1016/j.asoc.2018.04.039.
- [12] O. Cartagena, F. Trovò, M. Roveri, D. Sáez, Evolving fuzzy prediction intervals in nonstationary environments, *IEEE Transactions on Emerging Topics in Computational Intelligence* 8 (1) (2024) 903–916. doi:10.1109/TETCI.2023.3296486.
- [13] A. Khosravi, S. Nahavandi, D. Creighton, A. F. Atiya, Comprehensive review of neural network-based prediction intervals and new advances, *IEEE Transactions on Neural Networks* 22 (9) (2011) 1341–1352. doi:10.1109/TNN.2011.2162110.
- [14] N. Wang, W. Yao, Y. Zhao, X. Chen, Bayesian calibration of computer models based on Takagi–Sugeno fuzzy models, *Computer Methods in Applied Mechanics and Engineering* 378 (2021) 113724. doi:10.1016/j.cma.2021.113724.
- [15] X. Gong, J. Wang, Q. Ren, K. Zhang, E.-S. M. El-Alfy, J. Mańdziuk, Embedded feature selection approach based on TSK fuzzy system with sparse rule base for high-dimensional classification problems, *Knowledge-Based Systems* 295 (2024) 111809. doi:10.1016/j.knosys.2024.111809.
- [16] B. Zhang, J. Wang, Z. Shi, X. Gong, C. Zhang, K. Zhang, E.-S. M. El-Alfy, S. V. Ablameyko, First-order sparse TSK nonstationary fuzzy neural network based on the

- mean shift algorithm and the group lasso regularization, *Mathematics* 12 (1) (2023) 120.
doi:10.3390/math12010120.
- [17] D. Ji, Q. Fan, Q. Dong, Y. Liu, Convergence analysis of sparse TSK fuzzy systems based on spectral Dai-Yuan conjugate gradient and application to high-dimensional feature selection, *Neural Networks* 179 (2024) 106599. doi:10.1016/j.neunet.2024.106599.
- [18] A. E. Gelfand, D. K. Dey, Bayesian model choice: Asymptotics and exact calculations, *Journal of the Royal Statistical Society. Series B* 56 (3) (1994) 501–514. doi:10.1111/j.2517-6161.1994.tb01996.x.
- [19] H. Chipman, E. I. George, R. E. McCulloch, The practical implementation of Bayesian model selection, in: P. Lahiri (Ed.), *Model Selection*, Vol. 38 of *Institute of Mathematical Statistics Lecture Notes – Monograph Series*, Institute of Mathematical Statistics, 2001, pp. 65–116. doi:10.1214/lnms/1215540964.
- [20] M. Kasprzak, R. Giordano, T. Broderick, [How good is your Laplace approximation of the Bayesian posterior? finite-sample computable error bounds for a variety of useful divergences](#), *Journal of Machine Learning Research* 26 (2025) 1–81, paper 24-0619.
URL <https://jmlr.org/papers/volume26/24-0619/24-0619.pdf>
- [21] H. Ishwaran, J. S. Rao, Spike and slab variable selection: Frequentist and Bayesian strategies, *Annals of Statistics* 33 (2) (2005) 730–773. doi:10.1214/009053604000001147.
- [22] S. Jantre, S. Bhattacharya, T. Maiti, Spike-and-slab shrinkage priors for structurally sparse Bayesian neural networks, *IEEE Transactions on Neural Networks and Learning Systems* 36 (6) (2025) 11176–11188. doi:10.1109/TNNLS.2024.3485529.
- [23] Y. K. Huh, Y. Fu, K. Liu, A Bayesian spike-and-slab sensor selection approach for high-dimensional prognostics, *IEEE Transactions on Automation Science and Engineering* 22 (2025) 13814–13827. doi:10.1109/TASE.2025.3555165.
- [24] T. Denœux, Quantifying prediction uncertainty in regression using random fuzzy sets: The ENNreg model, *IEEE Transactions on Fuzzy Systems* 31 (10) (2023) 3690–3699. doi:10.1109/TFUZZ.2023.3268200.

- [25] H. Bian, J. Chang, G. Wang, N. R. Pal, Multihead Takagi–Sugeno–Kang fuzzy system for high-dimensional data, *IEEE Transactions on Fuzzy Systems* 33 (8) (2025) 2561–2573. doi:10.1109/TFUZZ.2025.3569227.
- [26] Z. Xue, Y. Xu, Y. Yang, S. Wang, DG-ALETSK: A high-dimensional fuzzy approach with simultaneous feature selection and rule extraction, *IEEE Transactions on Fuzzy Systems* 31 (11) (2023) 3852–3866. doi:10.1109/TFUZZ.2023.3270445.
- [27] V. Vovk, A. Gammerman, G. Shafer, *Algorithmic Learning in a Random World*, Springer, New York, NY, 2005. doi:10.1007/b106715.
- [28] J. Lei, A. Rinaldo, L. Wasserman, Distribution-free predictive inference for regression, *Journal of the American Statistical Association* 113 (523) (2018) 1094–1111. doi:10.1080/01621459.2017.1307116.
- [29] J. Pineau, P. Vincent, V. Gerstein, H. Larochelle, et al., *Improving reproducibility in machine learning research*, *Journal of Machine Learning Research* 22 (164) (2021) 1–20, a report from the NeurIPS 2019 Reproducibility Program. URL <https://jmlr.org/papers/volume22/20-303/20-303.pdf>
- [30] J. C. Bezdek, R. Ehrlich, W. Full, FCM: The fuzzy C-means clustering algorithm, *Computers & Geosciences* 10 (2–3) (1984) 191–203. doi:10.1016/0098-3004(84)90020-7.
- [31] J. A. Hoeting, D. Madigan, A. E. Raftery, C. T. Volinsky, Bayesian model averaging: A tutorial, *Statistical Science* 14 (4) (1999) 382–417. doi:10.1214/ss/1009212519.
- [32] T. M. Fragoso, W. Bertoli, F. Louzada, Bayesian model averaging: A systematic review and conceptual classification, *International Statistical Review* 86 (1) (2018) 1–28. doi:10.1111/insr.12243.
- [33] I.-C. Yeh, Modeling of strength of high-performance concrete using artificial neural networks, *Cement and Concrete Research* 28 (12) (1998) 1797–1808. doi:10.1016/S0008-8846(98)00165-3.
- [34] F. Wilcoxon, Individual comparisons by ranking methods, *Biometrics Bulletin* 1 (6) (1945) 80–83. doi:10.2307/3001968.

- 496 [35] C. E. Rasmussen, C. K. I. Williams, Gaussian Processes for Machine Learning, MIT Press,
497 Cambridge, MA, 2006.
- 498 [36] T. Gneiting, A. E. Raftery, Strictly proper scoring rules, prediction, and estimation, Jour-
499 nal of the American Statistical Association 102 (477) (2007) 359–378. [doi:10.1198/](https://doi.org/10.1198/016214506000001437)
500 [016214506000001437](https://doi.org/10.1198/016214506000001437).
- 501 [37] R. E. Kass, A. E. Raftery, Bayes factors, Journal of the American Statistical Association
502 90 (430) (1995) 773–795. [doi:10.1080/01621459.1995.10476572](https://doi.org/10.1080/01621459.1995.10476572).
- 503 [38] V. Ročková, E. I. George, The spike-and-slab LASSO, Journal of the American Statistical
504 Association 113 (521) (2018) 431–444. [doi:10.1080/01621459.2016.1260469](https://doi.org/10.1080/01621459.2016.1260469).
- 505 [39] C. M. Carvalho, N. G. Polson, J. G. Scott, The horseshoe estimator for sparse signals,
506 Biometrika 97 (2) (2010) 465–480. [doi:10.1093/biomet/asq017](https://doi.org/10.1093/biomet/asq017).